

StarCraft II 환경 소개

'StarCraft II로 배우는 강화학습' 웨비나

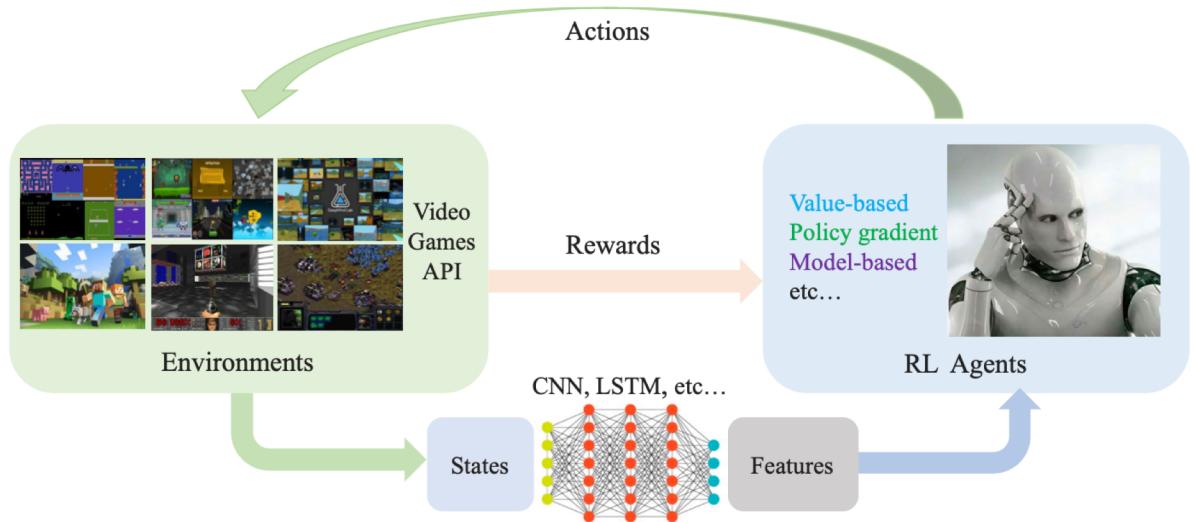
Jul 15, 2020

박석

DRL in Complex Gaming Environments

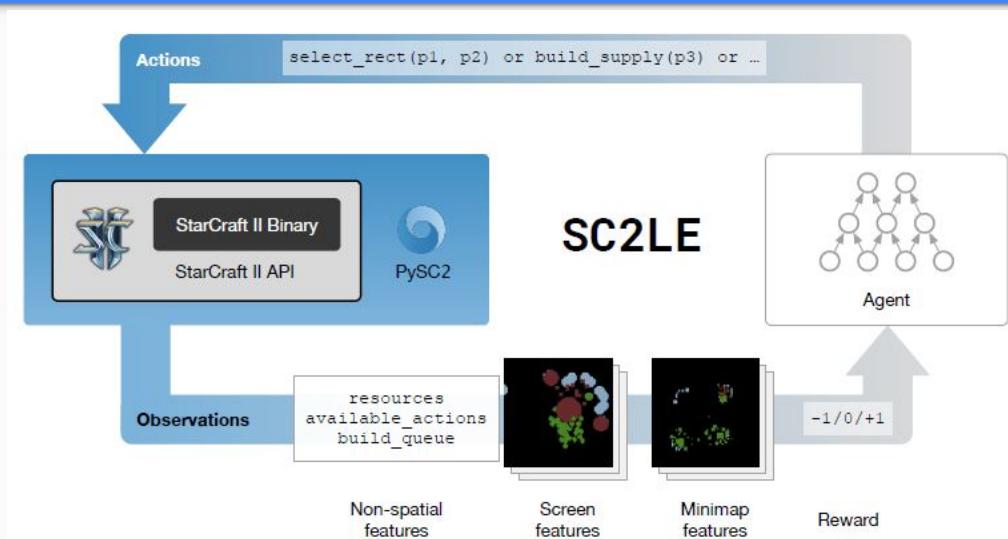
- 비디오 게임은 에이전트가 해결해야 할 흥미롭고 복잡한 문제를 제공하여 게임을 AI 연구에 완벽한 환경으로 만든다.
- 비디오 게임의 가상 환경은 안전하고 제어 가능하고, ML 알고리즘에 유용한 데이터를 무한대로 빠르게 제공한다.
- 병렬처리 가능.
- 가상 환경 제작에 소요되는 적은 비용.
- 평가의 객관성 확보.
- 이러한 특성은 **게임 환경을 DRL의 난제를 연구할 수 있는 주요 영역으로 만들었다.**

General DRL Gaming Environments Diagram



비디오 게임에 대한 일반적인 DRL의 프레임워크 다이어그램. 딥러닝 모델은 비디오 게임 API에서 입력을 받아 의미있는 피쳐를 자동으로 추출한다. DRL 에이전트는 이러한 피쳐를 기반으로 액션을 생성하고 환경을 다음 상태로 전환한다.

StarCraft II Environment Case – SC2LE



Ref : <https://deepmind.com/documents/110/sc2le.pdf>

What is PySC2?

- StarCraft II machine learning environment
- Interacts with Blizzard's StarCraft II API
- Python based
- Designed to emulate human abilities
- Backed by DeepMind
- github.com/deepmind/pysc2

What is Blizzard's StarCraft II API?

- Allows you to interact with the game via Protobuf
- Has limited Linux support
- Can play replays
- Provides the ability to investigate game state
- Provides the ability to perform player actions
- Supports 2 players
- Has some limitations
- Is still being developed

Ref : <https://blizzard.github.io/s2client-api/index.html>

Alternatives to PySC2

- Dave Churchill's CommandCenter - C++ framework for BW and SC2, very popular and used for bot battles
- Python SC2 - less human realistic
- C#, Clojure, Java, Go

Why PySC2?

- Python = TensorFlow, PyTorch and more
- Coded by DeepMind development team
- Possible to build a human-comparable bot like ‘AlphaStar’ by PYSC2 APIs

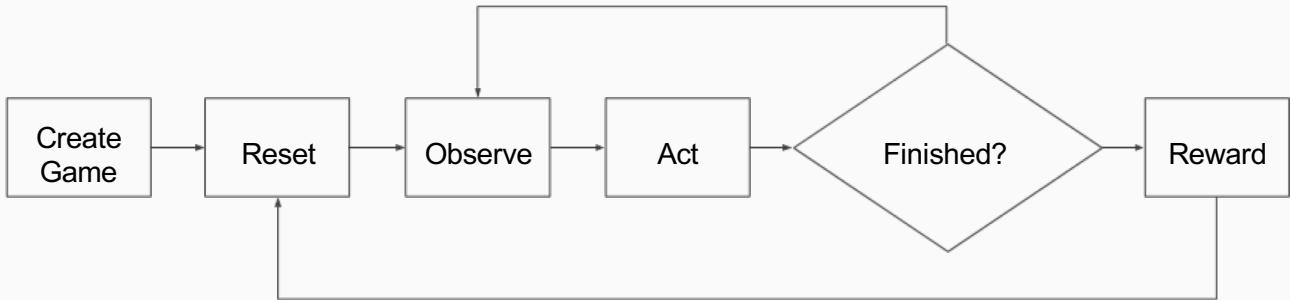
Why StarCraft II?

- Deterministic - there's no random chance, in the same state the same action will produce the same result
- Extremely Large State and Action Space - it takes a lot of time to explore and find what works and what doesn't
- Hidden Information - there is information about your opponent you don't know, exacerbated by fog of war

Sort of agents to build

- Completely scripted
- Partially scripted, partially ML
- Completely ML but structured specifically for SC2
- Completely ML with no SC2 specific structures

Game Engine Flow



Example Agent

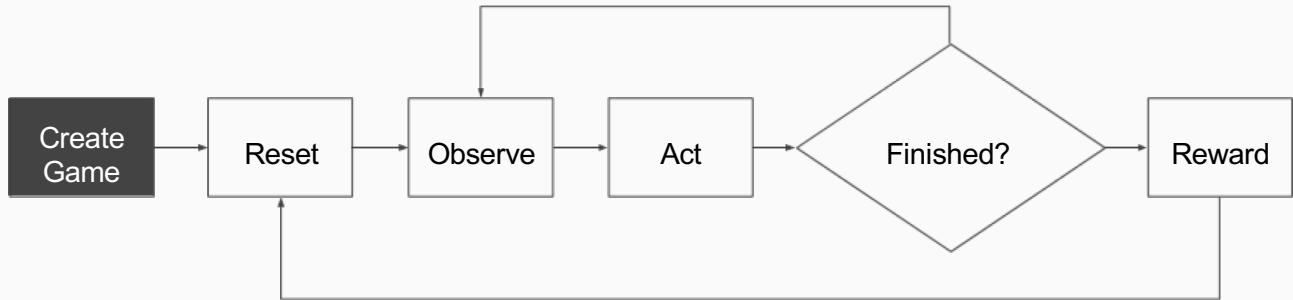
```
class MyAgent(base_agent.BaseAgent):
    def __init__(self):
        super(MyAgent, self).__init__()
        # One-time setup

    def reset(self):
        super(MyAgent, self).reset()
        # Before each game

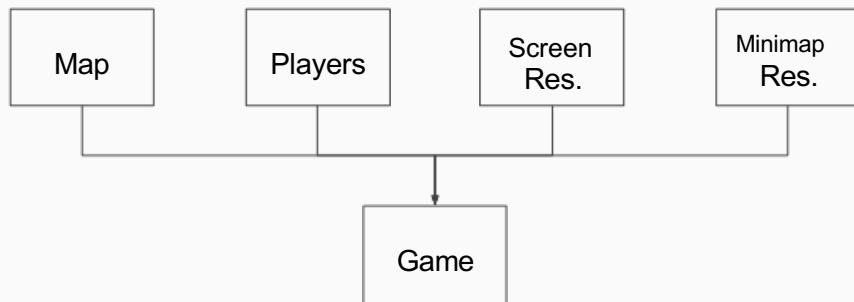
    def step(self, obs):
        super(MyAgent, self).step(obs)

        # Read state from obs and ALWAYS act
        return actions.FUNCTIONS.no_op()
```

Game Engine Flow – Create Game



Main Game Options



Code Example of Game Options

```
with sc2_env.SC2Env(  
    map_name="Simple64",  
    players=[sc2_env.Agent(sc2_env.Race.terran),  
            sc2_env.Bot(sc2_env.Race.zerg, sc2_env.Difficulty.easy)],  
    agent_interface_format=features.AgentInterfaceFormat(  
        feature_dimensions=features.Dimensions(screen=84,  
                                                minimap=64),  
        action_space=actions.ActionSpace.FEATURES),  
) as env:
```

Maps - Mini Games

Simplified games designed for testing algorithms

- BuildMarines
- CollectMineralsAndGas
- CollectMineralShards
- DefeatRoaches
- DefeatZerglingsAndBanelings
- FindAndDefeatZerglings
- MoveToBeacon



Maps - Melee Maps

- Designed for full game play
- Empty128
- Flat32/48/64/96/128 - No terrain variations
- Simple64/96/128 - Some terrain

Maps - Ladder Maps

- The usual ladder maps
- Can lag behind a bit, depending on Linux support and code updates
- Far more complex to generalise for due to variations in terrain, base locations, resources, etc.

Bot Difficulties

- very_easy
- easy
- medium
- medium_hard = Hard
- hard = Harder
- harder = Very hard
- very_hard = Elite
- cheat_vision
- cheat_money
- cheat_insane



Other Useful Game Options

step_mul

The number of steps to take before acting again, there are 22.4 steps per second for normal “faster” games, so a value of $8 = 168$ APM, I have had issues below a step mul of 2

game_steps_per_episode

The maximum steps to take before the game automatically ends

Other Useful Game Options

visualize

Shows a custom rendered version of the observations,
good for debugging but slows the game considerably

disable_fog

Disables the fog of war so everything is visible

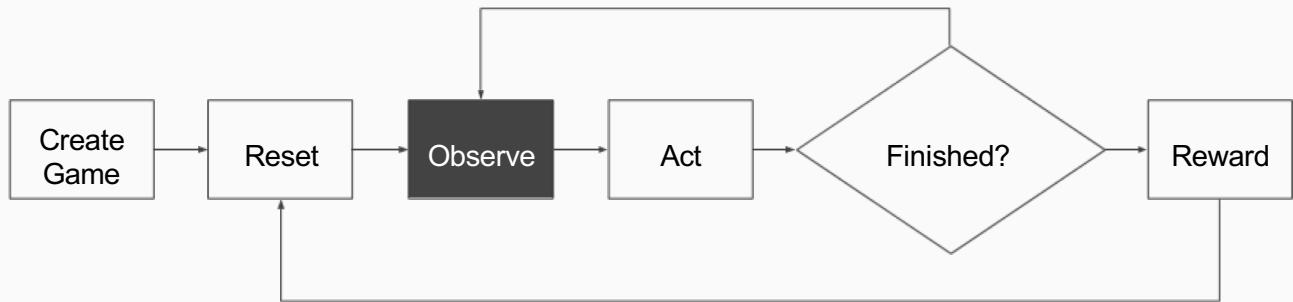
Machine Learning Tips

- Start against very easy bots
- Play as a single race
- Play against a single race
- Disable the fog of war
- Use a single, simple map
- Human players can beat very easy with an APM <60 so consider a step_mul of 20-30

Machine Learning Tips

- Limit the games to half an hour or so (40,320 steps)
- Consider having your agent play against a simple scripted bot
- Consider self-play, you get twice the learning per game but it may not generalise
- Compare your bot against one that chooses completely random actions

Game Engine Flow - Observe



Example State

```
state = (command_center_count,
          supply_depot_count,
          barracks_count,
          scv_count,
          marine_count,
          base1_enemy_count,
          base2_enemy_count,
          base3_enemy_count,
          base4_enemy_count,
          base1_friendly_count,
          base2_friendly_count,
          base3_friendly_count,
          base4_friendly_count)
```

Step “obs”

`obs.first()`

Whether or not this is the first step of the game, good for doing things like position and race detection

`obs.last()`

Whether or not this is the last step, use this to learn, save, and wrap things up

`obs.observation.reward`

Use this for sparse rewards, will be 0 for mid-steps or a draw, -1 for a loss, 1 for a win

Step “obs”

`obs.observation.game_loop`

The current step, you can work out the current second if you divide by 22.4

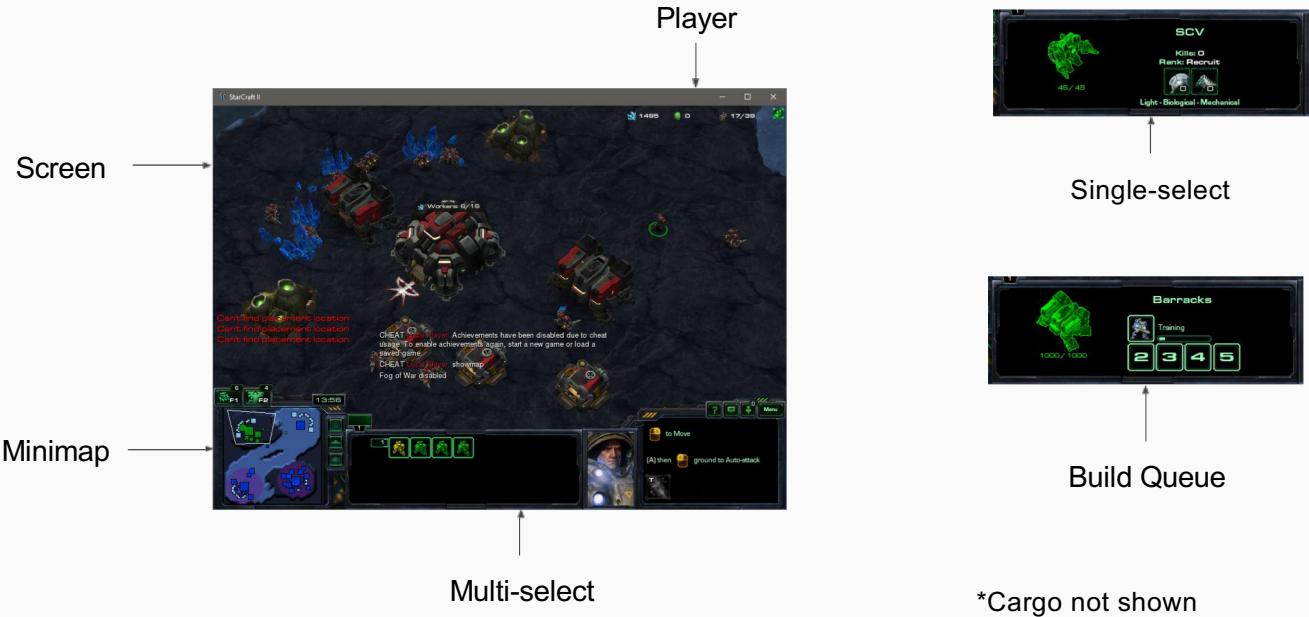
`obs.observation.available_actions`

Contains a list of actions that can be performed in the current state

`obs.observation`

Everything else

Main “Features”



Some Screen Features



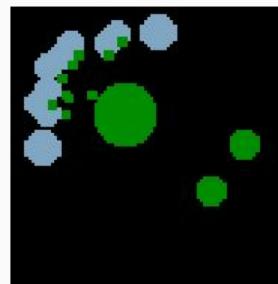
Height map



Unit type



Visibility



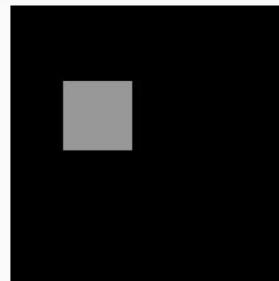
Player relative

Other important features include power, creep and effects

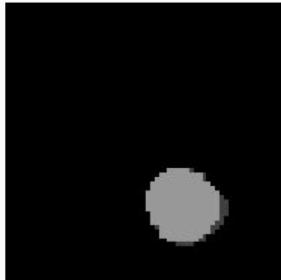
Some Minimap Features



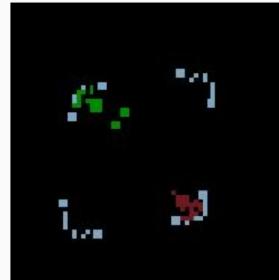
Height map



Camera location



Visibility



Player relative

Other important features include creep

Scalar Feature Layers

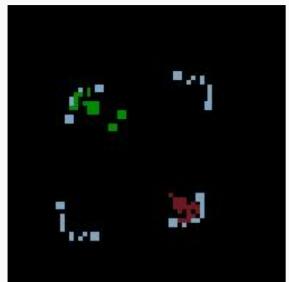
x	↓	
y →		
		[[0 0 0 0 0 0 0 0]
		[0 0 0 0 0 0 0 0]
		[0 255 255 255 212 212 0 0]
		[0 255 255 255 212 212 0 0]
		[0 0 212 0 212 0 0 0]
		[0 212 212 255 255 255 0 0]
		[0 212 212 212 255 255 0 0]
		[0 0 0 0 0 0 0 0]]



```
height_map = obs.observation.feature_minimap.height_map  
height_map[y][x]  
height_map[2][1] = 255
```

Categorical Feature Layers

```
[[0 0 0 0 0 0 0 0 0]
 [0 0 0 0 0 0 0 0 0]
 [0 3 1 0 3 3 0 0]
 [0 3 1 0 0 0 0 0]
 [0 0 0 0 0 0 0 0 0]
 [0 3 0 0 0 3 0 0]
 [0 3 3 0 3 3 0 0]
 [0 0 0 0 0 0 0 0 0]]
```

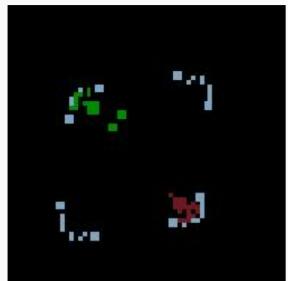


```
player_relative = obs.observation.feature_minimap.player_relative
```

Remember not to use ranges (e.g. 0-3) for categorical layers when feeding into neural networks, instead supply each category as 0 or 1.

Categorical Feature Layers

```
[[False False False False False False False False]
 [False False False False False False False False]
 [False False True False False False False False]
 [False False True False False False False False]
 [False False False False False False False False]]
```



```
player_relative = obs.observation.feature_minimap.player_relative
minimap_self = (player_relative == features.PlayerRelative.SELF)
```

Categorical Feature Layers

```
([2, 2], [2, 3]) # [y1, y2], [x1, x2]
```



```
player_relative = obs.observation.feature_minimap.player_relative  
player_y, player_x = (player_relative == features.PlayerRelative.SELF).nonzero()
```

Categorical Feature Layers

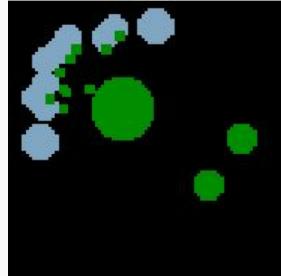
```
([2, 2], # [x1, y1]  
 [3, 2]) # [x2, y2]
```



```
player_relative = obs.observation.feature_minimap.player_relative  
player_y, player_x = (player_relative == features.PlayerRelative.SELF).nonzero()  
player_xy = zip(player_x, player_y)
```

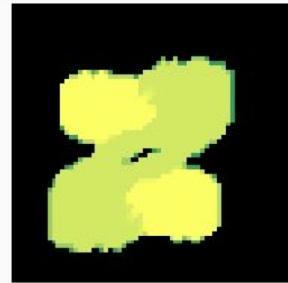
Screen Feature Notes

- The perspective is different, so don't expect everything to perfectly match the normal game
- Units can overlap and be difficult to identify



Minimap Feature Notes

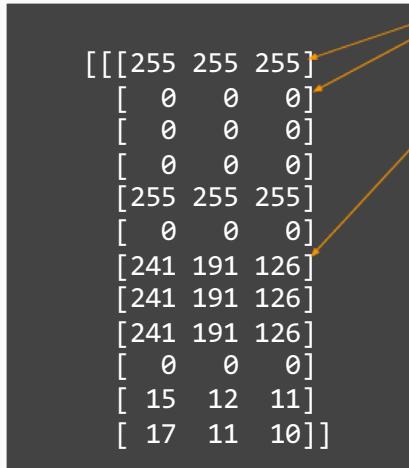
- The minimap in the API may not match the minimap in the game
- Each pixel on the minimap can only contain one detail per feature layer



Enable RGB Observations

```
with sc2_env.SC2Env(  
    map_name="Simple64",  
    players=[sc2_env.Agent(sc2_env.Race.terran),  
             sc2_env.Bot(sc2_env.Race.zerg, sc2_env.Difficulty.easy)],  
    agent_interface_format=features.AgentInterfaceFormat(  
        rgb_dimensions=features.Dimensions(screen=84,  
                                              minimap=64),  
        action_space=actions.ActionSpace.RGB),  
    ) as env:
```

RGB Observations



`obs.observation.rgb_minimap`
`obs.observation.rgb_screen`

*BGR instead of RGB may be a bug

RGB Feature Notes

- If you are observing in the RGB space you should act in the RGB space to maintain perspective
- Sizing does not seem to be exact (e.g. specifying 8x8 produced a 12x10 grid)
- Screen perspective matches the regular game
- Minimap seems to match the regular game

Single, Multi-Select and Cargo Observations

```
obs.observation.single_select  
obs.observation.multi_select  
obs.observation.cargo
```



```
[ [45  
  1  
  45  
  0  
  0  
  0  
  0]  
  Unit type (SCV)  
  Player relative (Self)  
  Health  
  Shields  
  Energy  
  Transport slots taken  
  Build progress (Not applicable to SCV, normally 0-100)
```

Player Observations

```
obs.observation.player.player_id  
obs.observation.player.minerals  
obs.observation.player.vespene  
obs.observation.player.food_used  
obs.observation.player.food_cap  
obs.observation.player.food_army  
obs.observation.player.food_workers  
obs.observation.player.idle_worker_count  
obs.observation.player.army_count  
obs.observation.player.warp_gate_count  
obs.observation.player.larva_count  
  
free_supply = food_cap - food_used
```

Enable Feature Units

```
with sc2_env.SC2Env(  
    map_name="Simple64",  
    players=[sc2_env.Agent(sc2_env.Race.terran),  
             sc2_env.Bot(sc2_env.Race.zerg, sc2_env.Difficulty.easy)],  
    agent_interface_format=features.AgentInterfaceFormat(  
        feature_dimensions=features.Dimensions(screen=84,  
                                                minimap=64),  
        action_space=actions.ActionSpace.FEATURES,  
        use_feature_units=True),  
) as env:
```

Feature Units

- Every visible unit on screen
- Exact unit location
- Build progress
- Assigned worker count
- Ideal worker count

```
marines = [unit for unit in obs.observation.feature_units
           if unit.unit_type == units.Terran.Marine
           and unit.alliance == features.PlayerRelative.SELF]
```

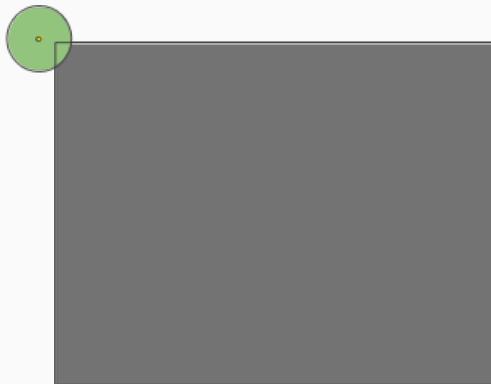
Feature Unit Properties

```
unit_type  
alliance  
health  
shield  
energy  
cargo_space_taken  
build_progress # 0-100  
health_ratio # 0-255  
shield_ratio # 0-255  
energy_ratio # 0-255  
display_type  
owner  
x  
y
```

```
facing  
radius  
cloak  
is_selected  
is_blink  
isPowered  
mineral_contents  
vespene_contents  
cargo_space_max  
assigned_harvesters  
ideal_harvesters  
weapon_cooldown  
order_length  
addon_unit_type # soon?
```

Feature Unit Notes

- Unit coordinates may be outside the screen since they are the centre of the unit, you will have to clip the values



Feature Unit Notes

- Unit visibility seems to match the real game, so if you are acting in the FEATURE space you may be able to move them to a location that makes them no longer visible

Machine Learning Tips

- Be sure to use categorical encoding for categorical feature layers, instead of ranges

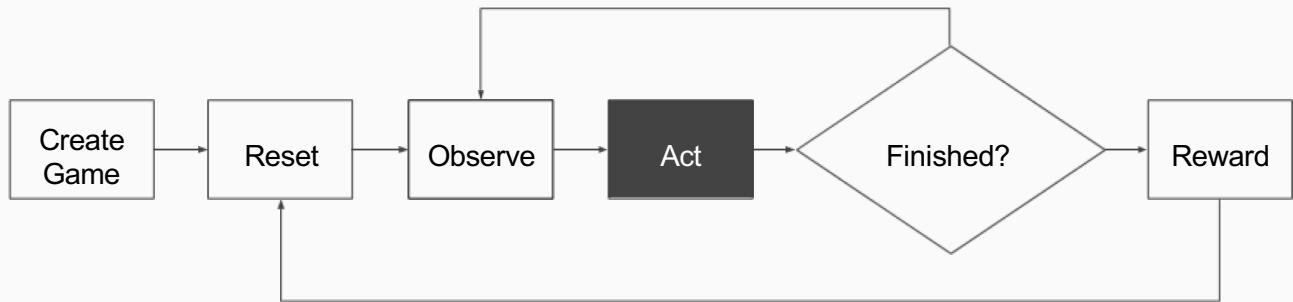
```
[0 1 0 2 2 0 3 0]  
[0 1 0 0 0 0 0 0]  
[0 0 0 1 1 0 0 0]  
[0 0 0 0 0 0 1 0]
```

```
0 1 2 3
```

Machine Learning Tips

- Feature scaling - things like minerals and step count can number into tens of thousands, while other values may be < 10
- Consider reducing minerals to “can afford” flags
- Crop and rotate the minimap so every game is from the same perspective

Game Engine Flow - Act



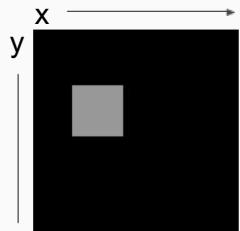
Move the Camera

Get the camera position:

```
camera_ys, camera_xs = (obs.observation.feature_minimap.camera == 1).nonzero()
camera_y = camera_ys.mean()
camera_x = camera_xs.mean()
camera_size = camera_y.max - camera_y.min # assuming square resolution
```

Set the camera position:

```
return actions.FUNCTIONS.move_camera((x, y)) # (x, y) is a tuple
```



- Due to rounding errors (or something) the camera position you receive may not match what you send in
- The x and y coordinates must be within the minimap size, e.g. 0-63
- Top-left is (0, 0)

Select an Idle Worker

```
if obs.observation.player.idle_worker_count > 0:  
    return actions.FUNCTIONS.select_idle_worker("select")
```

Selecting idle workers will move the screen

You can also “select_all”

Build a Barracks

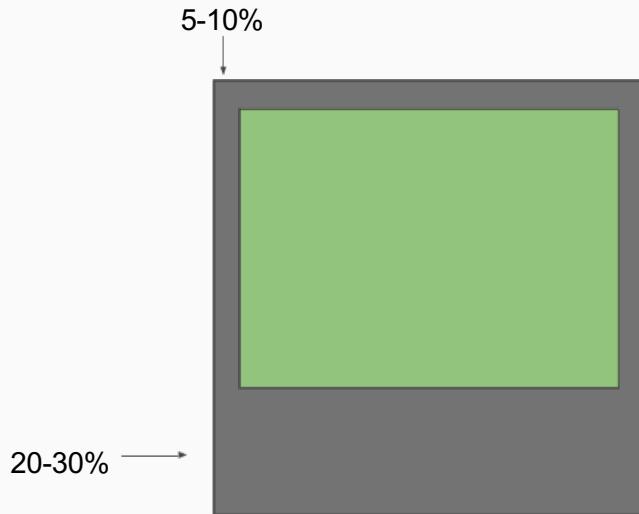
```
if (actions.FUNCTIONS.Build_Barracks_screen.id in  
    obs.observation.available_actions):  
    return actions.FUNCTIONS.Build_Barracks_screen("now", (x, y))
```

You can help your agent by preventing it from building in the mineral line

For Protoss or Zerg agents you can help by limiting to power or creep

Building Locations

It's best to add a margin to the left, top, and right, and a larger margin to the bottom



Autocast Repair

```
if (actions.FUNCTIONS.Effect_Repair_autocast.id in
    obs.observation.available_actions):
    return actions.FUNCTIONS.Effect_Repair_autocast()
```

Harvest Minerals

```
if (actions.FUNCTIONS.Harvest_Gather_screen.id in
    obs.observation.available_actions):
    minerals = [unit for unit in self.obs.observation.feature_units
                if unit.unit_type in MINERAL_UNIT_TYPES]
    if len(minerals) > 0:
        mineral = random.choice(minerals)

    return actions.FUNCTIONS.Harvest_Gather_screen(
        "queued", (mineral.x, mineral.y))
```

I have not had a lot of success with queued mineral harvesting

Mineral Unit Types

```
units.Neutral.BattleStationMineralField
units.Neutral.BattleStationMineralField750
units.Neutral.LabMineralField
units.Neutral.LabMineralField750
units.Neutral.MineralField
units.Neutral.MineralField750
units.Neutral.PurifierMineralField
units.Neutral.PurifierMineralField750
units.Neutral.PurifierRichMineralField
units.Neutral.PurifierRichMineralField750
units.Neutral.RichMineralField
units.Neutral.RichMineralField750
```

Vespene Unit Types

```
units.Neutral.ProtossVespeneGeyser  
units.Neutral.PurifierVespeneGeyser  
units.Neutral.RichVespeneGeyser  
units.Neutral.ShakurasVespeneGeyser  
units.Neutral.SpacePlatformGeyser  
units.Neutral.VespeneGeyser
```

Selecting all Barracks

```
barracks = [unit for unit in self.obs.observation.feature_units
            if unit.unit_type == units.Terran.Barracks
            and unit.alliance == features.PlayerRelative.SELF]
if len(barracks) > 0:
    barrack = random.choice(barracks)

return actions.FUNCTIONS.select_point(
    "select_all_type", (barrack.x, barrack.y))
```

Remember to clip the x and y coordinates to fit the screen resolution (e.g. 0-83)

Barrack is not the singular for barracks, I know

Set Rally Point

```
if (actions.FUNCTIONS.Rally_Units_screen.id in  
    obs.observation.available_actions):  
    return actions.FUNCTIONS.Rally_Units_screen("now", (x, y))
```

You can also rally workers

You can also use the minimap

Add Barracks to Control Group

```
return actions.FUNCTIONS.select_control_group("append", 0)
```

You can also set and recall

I have had issues with this

Train a Marine

```
if (actions.FUNCTIONS.Train_Marine_quick.id in  
    obs.observation.available_actions):  
    return actions.FUNCTIONS.Train_Marine_quick("now")
```

There seemed to be a bug currently that stops the correct distribution of build orders

In my experience when you have multiple units selected, all commands apply to all units even if normally they would not (e.g. SCVs and buildings)

Upgrade to Orbital Command

```
if (actions.FUNCTIONS.Morph_OrbitalCommand_quick.id in  
    obs.observation.available_actions):  
    return actions.FUNCTIONS.Morph_OrbitalCommand_quick("now")
```

Scan

```
if (actions.FUNCTIONS.Effect_Scan_screen.id in  
    obs.observation.available_actions):  
    return actions.FUNCTIONS.Effect_Scan_screen("now", (x, y))
```

You can also scan the minimap

Build Tech Lab

```
if (actions.FUNCTIONS.Build_TechLab_quick.id in  
    obs.observation.available_actions):  
    return actions.FUNCTIONS.Build_TechLab_quick("now")
```

There is currently a bug that only allows “Build_Techlab_screen”, so you need to lift with “Lift_Barracks_quick” and then use “Build_Techlab_screen” with coordinates

Action Grouping

- By default common actions are grouped, e.g. “burrow” for Zerg units has an individual command for each unit, but when they are grouped you only need to issue one command
- You can disable this by setting `hide_specific_actions=False` in your agent interface format

Research Stim

```
if (actions.FUNCTIONS.Research_Stimpack_quick.id in  
    obs.observation.available_actions):  
    return actions.FUNCTIONS.Research_Stimpack_quick("now")
```

It's hard to know if your research is done since research build queues don't exist (yet?), you may have to internally track how long it has been in progress, and check again if the research action is available (and you have enough resources)

Stim

```
if (actions.FUNCTIONS.Effect_Stim_quick.id in
    obs.observation.available_actions):
    return actions.FUNCTIONS.Effect_Stim_quick("now")
```

Attack

```
if (actions.FUNCTIONS.Attack_minimap.id in  
    obs.observation.available_actions):  
    return actions.FUNCTIONS.Attack_minimap("now", (x, y))
```

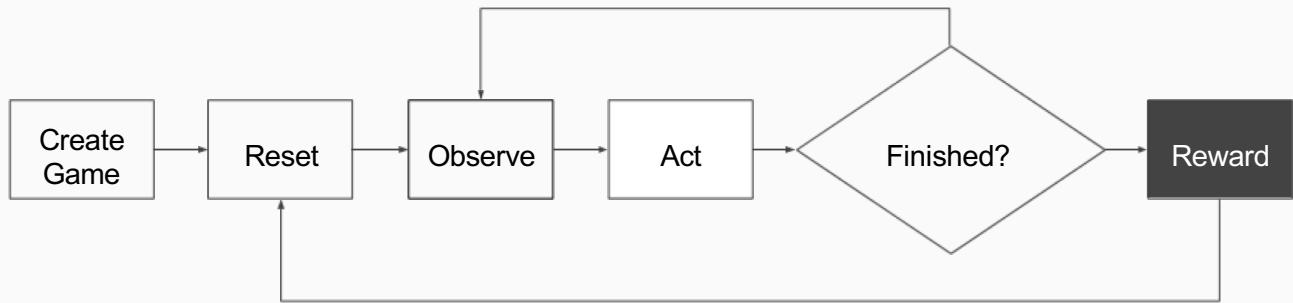
You can also attack the screen

You can also move units without attacking

Machine Learning Tips

- Consider scripted action sequences
- Consider scripted building locations
- Consider scripted build orders
- Attacking is simpler if you choose the closest enemy location
- Consider automated supply management
- Consider automated worker management

Game Engine Flow



PySC2 Rewards

```
if obs.last():
    reward = obs.reward

    # train your agent
    # save your learning

return actions.FUNCTIONS.no_op()
```

Machine Learning Tips

- You can define your own reward structure
- Be careful what you reward - you just might get it
- You might split your agent and reward for different things, for example:
 - Building units
 - Killing units
 - Finding the enemy

Limitations

- Cannot accept surrender
- Hard to track upgrades
- No research progress
- No ghosted buildings
- No build queue in multi-select
- Alerts/messages seem to be limited and unreliable

Windows vs Linux

- Linux is used by DeepMind, so that is most thoroughly tested
- Linux releases lag behind the latest balance changes and ladder maps
- Linux bugs take longer to get fixed
- Windows stays updated as it's the official game
- PySC2 often lags behind the latest Windows balance changes and ladder maps

More References

- Community : <https://discord.gg/b2gjyHR>
- DeepMind는 DeepRL 연구를 위해 StarCraft II를 사용하는 동기와 환경을 사용한 초기 연구 결과를 요약한 [블로그 포스트](#) 와 '[StarCraft II: A New Challenge for Reinforcement Learning](#)' [논문](#)을 발표했으니 참고.
- [Battle.net](#), [Liquipedia](#) 및 [Wikia](#)를 포함하여 Starcraft 온라인 자료.
- 맵 제작에 대해서는 [SC2Mapster](#)를 참조.